

# test

Levi Lee

11/21/2020

## Setup

First, we set working directories, install required libraries and import the data.

```
setwd("~/Documents/GitHub/Fall2020-Project4-group-4/doc")

packages.used <- c("dplyr", "ggplot2", "WeightedROC", "rpart", "rpart.plot")

# check packages that need to be installed.
packages.needed <- setdiff(packages.used, intersect(installed.packages()[,1], packages.used))

# install additional packages
if(length(packages.needed) > 0){
  install.packages(packages.needed, dependencies = TRUE)
}

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(WeightedROC)
library(rpart)
library(rpart.plot)

df_high <- read.csv("../data/highDim_dataset.csv")
df_low <- read.csv("../data/lowDim_dataset.csv")
```

## Introduction

*Description*

## About the Data

*Description*

## Background: Trees

*Description*

## Cross-Validation

*Description*

### Step 1: Set Controls and Establish Hyperparameters

*Description*

```
K <- 5 # number of CV folds
sample.reweight <- TRUE # run sample reweighting in model training

# setting the following to false loads data generated from a previous run
# this data is the same in each run due to a set seed

run.cv.trees_high <- FALSE # run cross-validation on the training set for trees on high dim data

run.cv.trees_low <- FALSE # run cross-validation on the training set for trees on low dim data
```

*Description*

```
# hyperparameters for trees
hyper_grid_trees <- expand.grid(
  cp = c(2(0), 2(-1), 2(-2), 2(-3), 2(-4),
        2(-5), 2(-6), 2(-7), 2(-8), 2(-9),
        2(-10), 2(-11), 2(-12), 2(-13), 2(-14),
        2(-15), 2(-16), 2(-17), 0, -2(0))
)
```

### Step 2: Cross-Validate the Hyperparameters

*Description*

*Description*

```
# features are the predictors: V1 - Vp
# column 1 is the response Y
# column 2 is the treatment A

feature_train_high = df_high[, -1:-2]
label_train_high = df_high[, 2]

feature_train_low = df_low[, -1:-2]
label_train_low = df_low[, 2]
```

## High Dimensional Data

*Description*

```
set.seed(5243)

if(run.cv.trees_high){
  res_cv_trees_high <- matrix(0, nrow = nrow(hyper_grid_trees), ncol = 4)
  for(i in 1:nrow(hyper_grid_trees)){
    cat("complexity = ", hyper_grid_trees$cp[i], "\n", sep = "")
  }
}
```

```

    res_cv_trees_high[i,] <- cv.function(features = feature_train_high, labels = label_train_high,
                                         cp = hyper_grid_trees$cp[i],
                                         K, reweight = sample.reweight)
    save(res_cv_trees_high, file = "../output/res_cv_trees_high.RData")
  }
} else {
  load("../output/res_cv_trees_high.RData")
}

```

## Low Dimensional Data

*Description*

```

set.seed(5243)

if(run.cv.trees_low){
  res_cv_trees_low <- matrix(0, nrow = nrow(hyper_grid_trees), ncol = 4)
  for(i in 1:nrow(hyper_grid_trees)){
    cat("complexity = ", hyper_grid_trees$cp[i], "\n", sep = "")
    res_cv_trees_low[i,] <- cv.function(features = feature_train_low, labels = label_train_low,
                                         cp = hyper_grid_trees$cp[i],
                                         K, reweight = sample.reweight)
    save(res_cv_trees_low, file="../output/res_cv_trees_low.RData")
  }
} else {
  load("../output/res_cv_trees_low.RData")
}

```

## Step 3: Visualize CV Error and AUC

*Description*

## High Dimensional Data

*Description*



```

# create data frame to organize results
res_cv_trees_high <- as.data.frame(res_cv_trees_high)
colnames(res_cv_trees_high) <- c("mean_error", "sd_error", "mean_AUC", "sd_AUC")
cv_results_trees_high = data.frame(hyper_grid_trees, res_cv_trees_high)

# look at top 5 models with highest AUC
cv_results_trees_high[order(cv_results_trees_high$mean_AUC, decreasing = TRUE), ][1:5, ]

##           cp mean_error  sd_error  mean_AUC    sd_AUC
## 8  7.812500e-03  0.4360668 0.03628091 0.5674343 0.03056252
## 6  3.125000e-02  0.4382665 0.02892494 0.5670732 0.03043522
## 7  1.562500e-02  0.4443627 0.01912652 0.5554088 0.01920727
## 5  6.250000e-02  0.4519064 0.01953579 0.5466031 0.02064263
## 18 7.629395e-06  0.4633120 0.01383644 0.5429351 0.02240466

```



```
best_cp_high <- cv_results_trees_high$cp[cv_results_trees_high$mean_AUC ==
                                         max(cv_results_trees_high$mean_AUC)]
```

```
best_cp_high
```

```
## [1] 0.0078125
```


## Low Dimensional Data


### Description

```
# create data frame to organize results
res_cv_trees_low <- as.data.frame(res_cv_trees_low)
colnames(res_cv_trees_low) <- c("mean_error", "sd_error", "mean_AUC", "sd_AUC")
cv_results_trees_low = data.frame(hyper_grid_trees, res_cv_trees_low)

# look at top 5 models with lowest AUC
cv_results_trees_low[order(cv_results_trees_low$mean_AUC, decreasing = TRUE), ][1:5, ]
```

```
##           cp mean_error  sd_error  mean_AUC    sd_AUC
## 18 7.629395e-06 0.3924126 0.07329778 0.6403787 0.08438962
## 3  2.500000e-01 0.3653965 0.06938198 0.6346035 0.06938198
## 4  1.250000e-01 0.3686625 0.05437925 0.6313375 0.05437925
## 5  6.250000e-02 0.3688654 0.03598451 0.6275346 0.03449488
## 6  3.125000e-02 0.3707763 0.04889355 0.6200507 0.04439815
```





```
best_cp_low <- cv_results_trees_low$cp[cv_results_trees_low$mean_AUC ==
                                     max(cv_results_trees_low$mean_AUC)]
```

```
best_cp_low
```

```
## [1] 7.629395e-06
```

## Propensity Score Estimation

*Description*

```
# imbalanced dataset requires weights
# to be used in the trained model
```

```
weights_high <- rep(NA, length(df_high$A))
for (v in unique(df_high$A)){
  weights_high[df_high$A == v] = 0.5 * length(df_high$A) / length(df_high$A[df_high$A == v])
}
```

```
weights_low <- rep(NA, length(df_low$A))
for (v in unique(df_low$A)){
  weights_low[df_low$A == v] = 0.5 * length(df_low$A) / length(df_low$A[df_low$A == v])
}
```

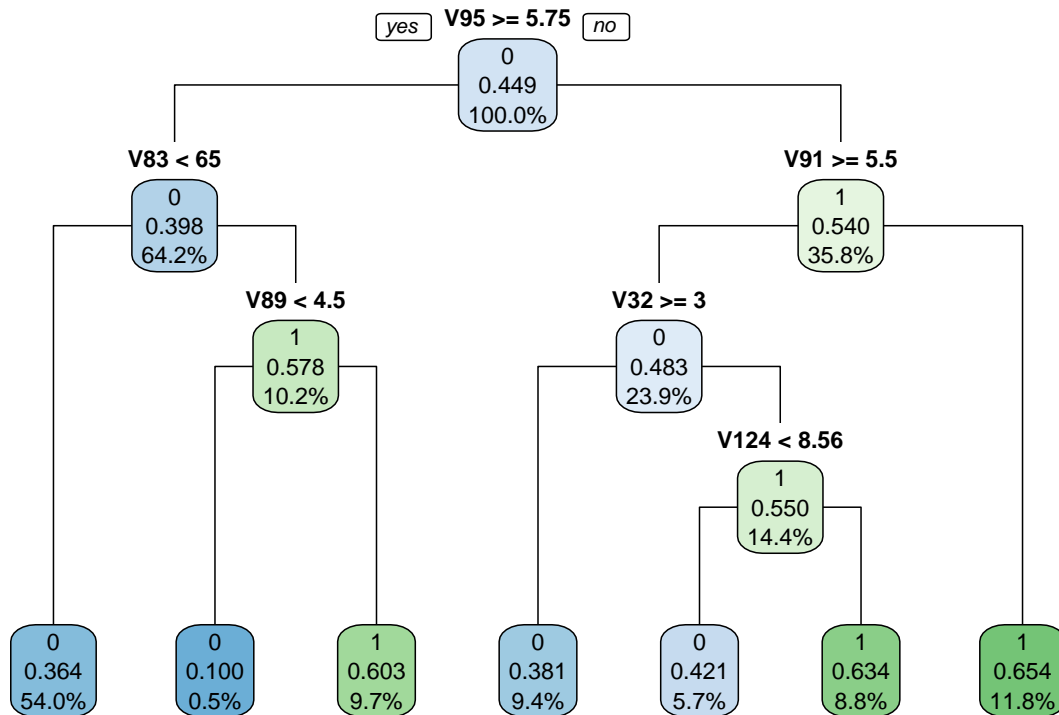
## High Dimensional Data

```
##### SET START TIME HERE
```

```
# create tree model for high dimensional data with best cp parameter
tree_high <- rpart(A ~ . - Y, method = "class", data = df_high, cp = best_cp_high)
```

```
# calculate propensity scores
prop_score_high <- predict(tree_high, newdata = df_high[, -2], type = "prob")[, 2]

##### SET END TIME HERE
```



## Low Dimensional Data

```
##### SET START TIME HERE

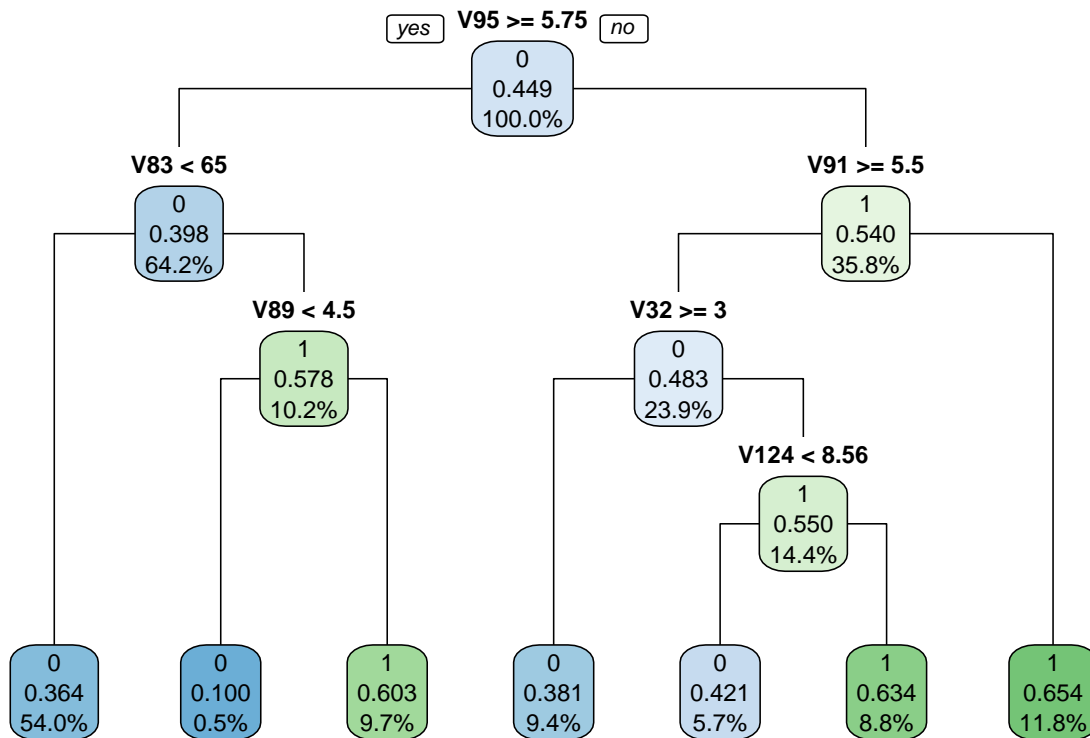
# create tree model for low dimensional data with best cp parameter
tree_low <- rpart(A ~ . - Y, method = "class", data = df_low, cp = best_cp_low)

# calculate propensity scores
prop_score_low <- predict(tree_low, newdata = df_low[, -2], type = "prob")[, 2]

##### SET END TIME HERE
```







```

ps_RA_high <- predict(regression_adjustment_high, df_high, type = "prob")
high_data_ps <- cbind(ps_RA_high, df_high)
pred_high <- lm(Y ~ A + ps_RA_high, data = high_data_ps)
summary(pred_high)

```

```

##
## Call:
## lm(formula = Y ~ A + ps_RA_high, data = high_data_ps)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.4713  -3.4878  -0.6694   2.7522  30.0897
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.9564     0.6207   9.596  <2e-16 ***
## A             -2.5271     0.2569  -9.836  <2e-16 ***
## ps_RA_high0 -30.5718     1.0322 -29.617  <2e-16 ***
## ps_RA_high1      NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.535 on 1997 degrees of freedom
## Multiple R-squared:  0.3068, Adjusted R-squared:  0.3061
## F-statistic: 441.8 on 2 and 1997 DF,  p-value: < 2.2e-16
ATE_RA_high = pred_high$coefficients[2]
ATE_RA_high

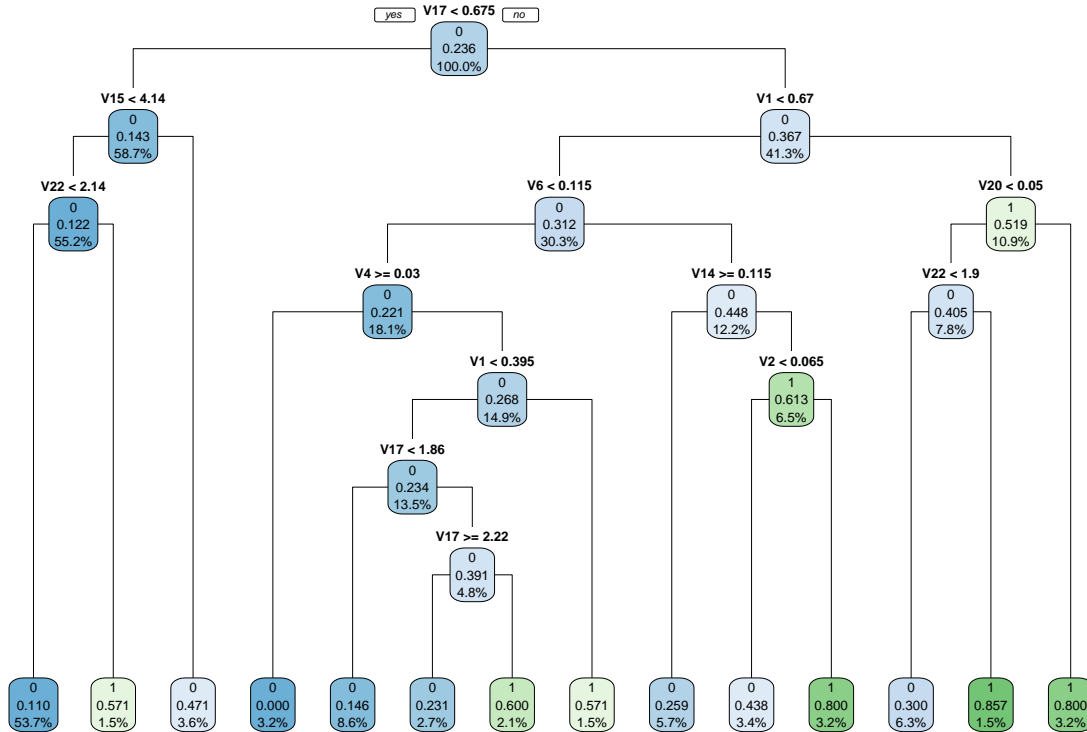
##           A
## -2.527116

```

## Low Dimensional Data

Description Also need run times here

```
regression_adjustment_low <- rpart(A ~ . - Y, data = df_low, method = "class", cp = 7.629395e-06)
rpart.plot(regression_adjustment_low, type = 1, digits = 3, fallen.leaves = TRUE)
```



```
ps_RA_low <- predict(regression_adjustment_low, df_low, type = "prob")
low_data_ps <- cbind(ps_RA_low, df_low)
pred_low <- lm(Y ~ A + ps_RA_low, data = low_data_ps)
summary(pred_low)
```

```
##
## Call:
## lm(formula = Y ~ A + ps_RA_low, data = low_data_ps)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.5153  -2.6807  -0.5655   1.8184  26.9282
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  22.5693     0.8535  26.445 < 2e-16 ***
## A             3.0532     0.5089   5.999 3.95e-09 ***
## ps_RA_low0   -7.5200     1.0017  -7.507 3.04e-13 ***
## ps_RA_low1      NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.055 on 472 degrees of freedom
## Multiple R-squared:  0.2829, Adjusted R-squared:  0.2798
## F-statistic: 93.08 on 2 and 472 DF, p-value: < 2.2e-16
```

```
ATE_RA_low = pred_low$coefficients[2]  
ATE_RA_low
```

```
##           A  
## 3.05324
```

## Stratification and Regression Adjustment

*Description*

## High Dimensional Data

*Description Also need run times here*

## Low Dimensional Data

*Description Also need run times here*

## Results

*Insert Comparison of ATE and all Runtimes Here*

## Conclusion

*Description*

## References

*Description*