

# Project 4: Causal Inference Algorithms Evaluation

Group 4: Zhenglei Chen, Jaival Desai, Qinzhe Hu, Levi Lee, Luyao Sun, Xinyi Wei

12/02/2020

## Setup

First, we set working directories as needed, install required libraries and import the data.

```
df_high <- read.csv("../data/highDim_dataset.csv")
df_low <- read.csv("../data/lowDim_dataset.csv")
```

## Introduction

### About the Data

```
## [1] "High Dimensional Data"
```

```
##
##      0      1
## 0.5515 0.4485
```

```
## [1] "Low Dimensional Data"
```

```
##
##      0      1
## 0.7642105 0.2357895
```

## Background on Classification/Regression Trees

### Cross-Validation

#### Step 1: Set Controls and Establish Hyperparameters

```
K <- 5 # number of CV folds
sample.reweight <- TRUE # run sample reweighting in model training

# setting the following to false loads data generated from a previous run
# this data is the same in each run due to a set seed

run.cv.trees_high <- FALSE # run cross-validation on the training set for trees on high dim data
run.cv.trees_low <- FALSE # run cross-validation on the training set for trees on low dim data
```

```
# hyperparameters for trees
hyper_grid_trees <- expand.grid(
  cp = c(20, 2-1, 2-2, 2-3, 2-4,
        2-5, 2-6, 2-7, 2-8, 2-9,
        2-10, 2-11, 2-12, 2-13, 2-14,
        2-15, 2-16, 2-17, 0, -20)
)
```

## Step 2: Cross-Validate the Hyperparameters

```
# features are the predictors: V1 - Vp
# column 1 is the response Y
# column 2 is the treatment A

feature_train_high = df_high[, -1:-2]
label_train_high = df_high[, 2]

feature_train_low = df_low[, -1:-2]
label_train_low = df_low[, 2]
```

```
set.seed(5243)

if(run.cv.trees_high){
  res_cv_trees_high <- matrix(0, nrow = nrow(hyper_grid_trees), ncol = 4)
  for(i in 1:nrow(hyper_grid_trees)){
    cat("complexity = ", hyper_grid_trees$cp[i], "\n", sep = "")
    res_cv_trees_high[i,] <- cv.function(features = feature_train_high,
                                         labels = label_train_high,
                                         cp = hyper_grid_trees$cp[i],
                                         K, reweight = sample.reweight)
    save(res_cv_trees_high, file = "../output/res_cv_trees_high.RData")
  }
} else{
  load("../output/res_cv_trees_high.RData")
}
```

## High Dimensional Data

```
set.seed(5243)

if(run.cv.trees_low){
  res_cv_trees_low <- matrix(0, nrow = nrow(hyper_grid_trees), ncol = 4)
  for(i in 1:nrow(hyper_grid_trees)){
    cat("complexity = ", hyper_grid_trees$cp[i], "\n", sep = "")
    res_cv_trees_low[i,] <- cv.function(features = feature_train_low,
                                         labels = label_train_low,
```

```

                                cp = hyper_grid_trees$cp[i],
                                K, reweight = sample.reweight)
  save(res_cv_trees_low, file="../output/res_cv_trees_low.RData")
}
}else{
  load("../output/res_cv_trees_low.RData")
}

```

## Low Dimensional Data

### Step 3: Visualize CV Error and AUC

```

# create data frame to organize results
res_cv_trees_high <- as.data.frame(res_cv_trees_high)
colnames(res_cv_trees_high) <- c("mean_error", "sd_error", "mean_AUC", "sd_AUC")
cv_results_trees_high = data.frame(hyper_grid_trees, res_cv_trees_high)

# look at top 5 models with highest AUC
cv_results_trees_high[order(cv_results_trees_high$mean_AUC, decreasing = TRUE), ][1:5, ]



```

## High Dimensional Data

```

##           cp mean_error  sd_error  mean_AUC    sd_AUC
## 8  7.812500e-03  0.4360668 0.03628091 0.5674343 0.03056252
## 6  3.125000e-02  0.4382665 0.02892494 0.5670732 0.03043522
## 7  1.562500e-02  0.4443627 0.01912652 0.5554088 0.01920727
## 5  6.250000e-02  0.4519064 0.01953579 0.5466031 0.02064263
## 18 7.629395e-06  0.4633120 0.01383644 0.5429351 0.02240466

```



```
best_cp_high <- cv_results_trees_high$cp[cv_results_trees_high$mean_AUC ==
                                         max(cv_results_trees_high$mean_AUC)]
```

```
best_cp_high
```

```
## [1] 0.0078125
```

```


# create data frame to organize results
res_cv_trees_low <- as.data.frame(res_cv_trees_low)
colnames(res_cv_trees_low) <- c("mean_error", "sd_error", "mean_AUC", "sd_AUC")
cv_results_trees_low = data.frame(hyper_grid_trees, res_cv_trees_low)

# look at top 5 models with lowest AUC
cv_results_trees_low[order(cv_results_trees_low$mean_AUC, decreasing = TRUE), ][1:5, ]


```

## Low Dimensional Data

##	cp	mean_error	sd_error	mean_AUC	sd_AUC
## 18	7.629395e-06	0.3924126	0.07329778	0.6403787	0.08438962
## 3	2.500000e-01	0.3653965	0.06938198	0.6346035	0.06938198
## 4	1.250000e-01	0.3686625	0.05437925	0.6313375	0.05437925
## 5	6.250000e-02	0.3688654	0.03598451	0.6275346	0.03449488
## 6	3.125000e-02	0.3707763	0.04889355	0.6200507	0.04439815



Mean of AUC with Different cp Values for Low Dimensional Data



```
best_cp_low <- cv_results_trees_low$cp[cv_results_trees_low$mean_AUC ==
                                     max(cv_results_trees_low$mean_AUC)]
```

```
best_cp_low
```

```
## [1] 7.629395e-06
```

## Propensity Score Estimation

```
# imbalanced dataset requires weights
# to be used in the trained model
```

```
weights_high <- rep(NA, length(df_high$A))
for (v in unique(df_high$A)){
  weights_high[df_high$A == v] = 0.5 * length(df_high$A) / length(df_high$A[df_high$A == v])
}
```

```
weights_low <- rep(NA, length(df_low$A))
for (v in unique(df_low$A)){
  weights_low[df_low$A == v] = 0.5 * length(df_low$A) / length(df_low$A[df_low$A == v])
}
```

```
start.time_propensity_score_high <- Sys.time()
```

```

# create tree model for high dimensional data with best cp parameter
tree_high <- rpart(A ~ . - Y, method = "class", data = df_high, cp = best_cp_high)


# calculate propensity scores
prop_score_high <- predict(tree_high, newdata = df_high[, -2], type = "prob")[, 2]

end.time_propensity_score_high <- Sys.time()
time_propensity_score_high <- end.time_propensity_score_high - start.time_propensity_score_high
time_propensity_score_high

```

## High Dimensional Data

## Time difference of 1.243244 secs



```

start.time_propensity_score_low <- Sys.time()

# create tree model for low dimensional data with best cp parameter
tree_low <- rpart(A ~ . - Y, method = "class", data = df_low, cp = best_cp_low)

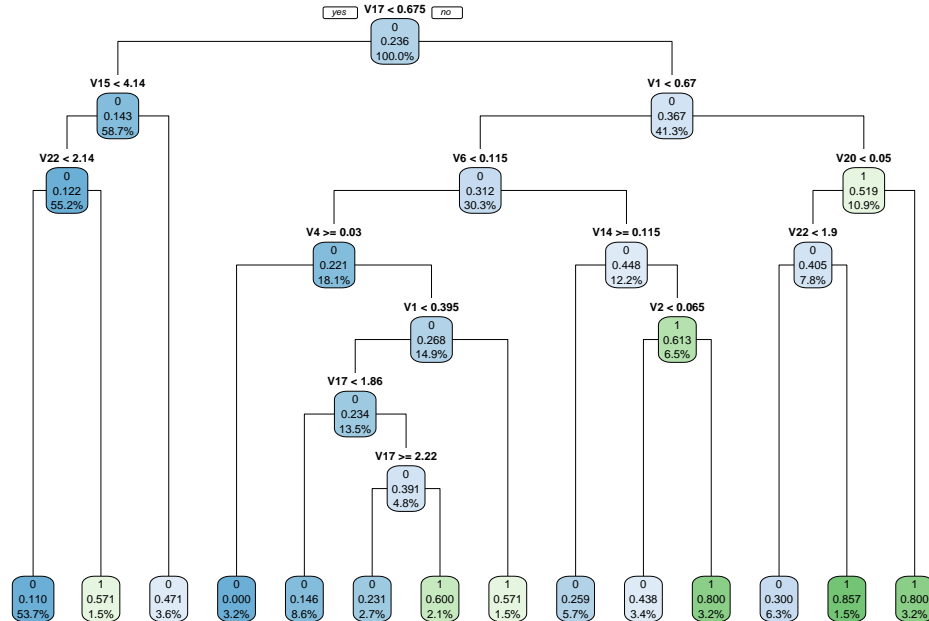
# calculate propensity scores
prop_score_low <- predict(tree_low, newdata = df_low[, -2], type = "prob")[, 2]

end.time_propensity_score_low <- Sys.time()
time_propensity_score_low <- end.time_propensity_score_low - start.time_propensity_score_low
time_propensity_score_low

```

## Low Dimensional Data

## Time difference of 0.03291011 secs



## ATE Estimation

With the estimated propensity scores on hand, we propose, explain, and discuss the pros and cons of three different ATE estimation algorithms: stratification, regression adjustment, and stratification plus regression adjustment together.

### Stratification

```
K = 3
strata <- seq(0, 1, by = 1/K)
```

```
start.time_stratification_high <- Sys.time()

df_high <- cbind(df_high, prop_score_high)
stratum_values_high <- rep(NA, length(strata))

for (i in 1:length(strata)){
  stratum_values_high[i] <- quantile(prop_score_high, strata[i])
}
```



```
# values of strata for high data
stratum_values_high
```

## High Dimensional Data

```
## [1] 0.1000000 0.3635523 0.4210526 0.6540084
```

```
df_high$stratum_class_high <- rep(NA, nrow(df_high))

# assign stratum class to each observation
for (i in 1:nrow(df_high)){
  if ((stratum_values_high[1] <= df_high$prop_score_high[i]) &
      (df_high$prop_score_high[i] < stratum_values_high[2])) {
    df_high$stratum_class_high[i] <- 1
  } else if ((stratum_values_high[2] <= df_high$prop_score_high[i]) &
             (df_high$prop_score_high[i] < stratum_values_high[3])) {
    df_high$stratum_class_high[i] <- 2
  } else if ((stratum_values_high[3] <= df_high$prop_score_high[i]) &
             (df_high$prop_score_high[i] <= stratum_values_high[4])) {
    df_high$stratum_class_high[i] <- 3
  }
}

summary_high = expand.grid(
  A = c(0, 1),
  stratum = seq(1, K, by = 1),
  n = NA,
  prop = NA,
  avg_y = NA
)

for (i in 1:nrow(summary_high)) {
  subset <- df_high[(df_high$A == summary_high$A[i]) &
                    (df_high$stratum_class_high == summary_high$stratum[i]), ]
  summary_high$n[i] = nrow(subset)
  summary_high$prop[i] = summary_high$n[i]/nrow(df_high)
  summary_high$avg_y[i] = mean(subset$Y)
}

for (i in 1:nrow(summary_high)) {
  if (is.nan(summary_high$avg_y[i]) == TRUE) {
    summary_high$avg_y[i] <- 0
  }
}

# this table records the mean response in each stratum; needed for stratification
summary_high
```

```
##   A stratum    n   prop    avg_y
## 1 0         1   9 0.0045 -1.556754
## 2 1         1   1 0.0005  3.448809
## 3 0        2 805 0.4025 -13.679379
```

```
## 4 1      2 465 0.2325 -16.168680
## 5 0      3 289 0.1445  -6.706643
## 6 1      3 431 0.2155  -8.342732
```

```
stratum_prop_high <- summary_high %>% group_by(stratum) %>% summarise(sum = sum(n)/nrow(df_high))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
# this table records the proportions for each stratum; also needed for stratification
stratum_prop_high
```

```
## # A tibble: 3 x 2
##   stratum  sum
##   <dbl> <dbl>
## 1      1  0.005
## 2      2  0.635
## 3      3  0.36
```

```
ATE_stratification_high = stratum_prop_high$sum[1]*(summary_high$avg_y[2] - summary_high$avg_y[1]) +
  stratum_prop_high$sum[2]*(summary_high$avg_y[4] - summary_high$avg_y[3]) +
  stratum_prop_high$sum[3]*(summary_high$avg_y[6] - summary_high$avg_y[5])
```

```
ATE_stratification_high
```

```
## [1] -2.14467
```

```
end.time_stratification_high <- Sys.time()
time_stratification_high <- end.time_stratification_high - start.time_stratification_high
time_stratification_high
```

```
## Time difference of 0.3165069 secs
```

We find that the ATE for the high dimensional dataset was -2.145 with a runtime of 0.317 seconds.

```
start.time_stratification_low <- Sys.time()

df_low <- cbind(df_low, prop_score_low)
stratum_values_low <- rep(NA, length(strata))

for (i in 1:length(strata)){
  stratum_values_low[i] <- quantile(prop_score_low, strata[i])
}

# values of strata for low data
stratum_values_low
```

## Low Dimensional Data

```
## [1] 0.0000000 0.1098039 0.2307692 0.8571429
```

```

df_low$stratum_class_low <- rep(NA, nrow(df_low))

# assign stratum class to each observation
for (i in 1:nrow(df_low)){
  if ((stratum_values_low[1] <= df_low$prop_score_low[i]) &
      (df_low$prop_score_low[i] < stratum_values_low[2])) {
    df_low$stratum_class_low[i] <- 1
  } else if ((stratum_values_low[2] <= df_low$prop_score_low[i]) &
              (df_low$prop_score_low[i] < stratum_values_low[3])) {
    df_low$stratum_class_low[i] <- 2
  } else if ((stratum_values_low[3] <= df_low$prop_score_low[i]) &
              (df_low$prop_score_low[i] <= stratum_values_low[4])) {
    df_low$stratum_class_low[i] <- 3
  }
}

summary_low = expand.grid(
  A = c(0, 1),
  stratum = seq(1, K, by = 1),
  n = NA,
  prop = NA,
  avg_y = NA
)

for (i in 1:nrow(summary_low)) {
  subset <- df_low[(df_low$A == summary_low$A[i]) &
                    (df_low$stratum_class_low == summary_low$stratum[i]), ]
  summary_low$n[i] = nrow(subset)
  summary_low$prop[i] = summary_low$n[i]/nrow(df_low)
  summary_low$avg_y[i] = mean(subset$Y)
}

for (i in 1:nrow(summary_low)) {
  if (is.nan(summary_low$avg_y[i]) == TRUE) {
    summary_low$avg_y[i] <- 0
  }
}

# this table records the mean response in each stratum; needed for stratification
summary_low

```

```

##   A stratum   n      prop    avg_y
## 1 0         1 15 0.03157895 18.25654
## 2 1         1  0 0.00000000  0.00000
## 3 0         2 262 0.55157895 15.34302
## 4 1         2  34 0.07157895 18.68881
## 5 0         3  86 0.18105263 19.14737
## 6 1         3  78 0.16421053 22.52296

```

```

stratum_prop_low <- summary_low %>% group_by(stratum) %>% summarise(sum = sum(n)/nrow(df_low))

## `summarise()` ungrouping output (override with `.groups` argument)

# this table records the proportions for each stratum; also needed for stratification
stratum_prop_low

## # A tibble: 3 x 2
##   stratum    sum
##   <dbl>  <dbl>
## 1      1  0.0316
## 2      2  0.623
## 3      3  0.345

ATE_stratification_low = stratum_prop_low$sum[1]*(summary_low$avg_y[2] - summary_low$avg_y[1]) +
  stratum_prop_low$sum[2]*(summary_low$avg_y[4] - summary_low$avg_y[3]) +
  stratum_prop_low$sum[3]*(summary_low$avg_y[6] - summary_low$avg_y[5])

ATE_stratification_low

## [1] 2.673899

end.time_stratification_low <- Sys.time()
time_stratification_low <- end.time_stratification_low - start.time_stratification_low
time_stratification_low

## Time difference of 0.121556 secs

```

We find that the ATE for the low dimensional dataset was 2.674 with a runtime of 0.122 seconds.

## Regression Adjustment

In this method, we regress the response variable ( $Y$ ) with the treatment variable ( $A$ ) and the propensity scores estimated using our model above, in this case, trees. The estimated coefficient of the treatment variable ( $A$ ) is then an estimate of the ATE.

D'Agostino (1998) and Austin (2011) compare regression adjustment with more traditional propensity score methods. One of the main advantages of the regression adjustment is in its simplicity in execution, in which one performs a somewhat basic linear regression model on two covariates and one response variable.

However, depending on the size of the dataset, this may run into computation issues as linear regression involves finding the inverse of a matrix. Additionally, regression adjustment may also not be helpful in cases where there is a strong separation between the two groups.

No such issues were present in this setup given that both datasets had a relatively small number of observations and there is no clear separation between the two groups, as shown in the residual plots below.

```

start.time_regression_adjustment_high <- Sys.time()

ps_RA_high <- predict(tree_high, df_high, type = "prob")
high_data_ps <- cbind(ps_RA_high, df_high)
pred_high <- lm(Y ~ A + ps_RA_high, data = high_data_ps)
summary(pred_high)

```

## High Dimensional Data

```

##
## Call:
## lm(formula = Y ~ A + ps_RA_high, data = high_data_ps)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.4713  -3.4878  -0.6694   2.7522  30.0897
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.9564     0.6207   9.596 <2e-16 ***
## A            -2.5271     0.2569  -9.836 <2e-16 ***
## ps_RA_high0 -30.5718     1.0322 -29.617 <2e-16 ***
## ps_RA_high1      NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.535 on 1997 degrees of freedom
## Multiple R-squared:  0.3068, Adjusted R-squared:  0.3061
## F-statistic: 441.8 on 2 and 1997 DF,  p-value: < 2.2e-16

```

```

ATE_regression_adjustment_high = pred_high$coefficients[2]
ATE_regression_adjustment_high

```

```

##      A
## -2.527116

```

```

end.time_regression_adjustment_high <- Sys.time()
time_regression_adjustment_high <- end.time_regression_adjustment_high -
  start.time_regression_adjustment_high
time_regression_adjustment_high


```

```

## Time difference of 0.0410831 secs

```

## Residual Plot of Regression Adjustment Model – High Dim



We find that the ATE for the high dimensional dataset was -2.527 with a runtime of 0.041 seconds.

```
start.time_regression_adjustment_low <- Sys.time()

ps_RA_low <- predict(tree_low, df_low, type = "prob")
low_data_ps <- cbind(ps_RA_low, df_low)
pred_low <- lm(Y ~ A + ps_RA_low, data = low_data_ps)
summary(pred_low)
```

## Low Dimensional Data

```
##
## Call:
## lm(formula = Y ~ A + ps_RA_low, data = low_data_ps)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.5153  -2.6807  -0.5655   1.8184  26.9282
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  22.5693     0.8535  26.445 < 2e-16 ***
## A             3.0532     0.5089   5.999 3.95e-09 ***
## ps_RA_low0   -7.5200     1.0017  -7.507 3.04e-13 ***
## ps_RA_low1      NA           NA      NA      NA
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.055 on 472 degrees of freedom
## Multiple R-squared:  0.2829, Adjusted R-squared:  0.2798
## F-statistic: 93.08 on 2 and 472 DF,  p-value: < 2.2e-16
```

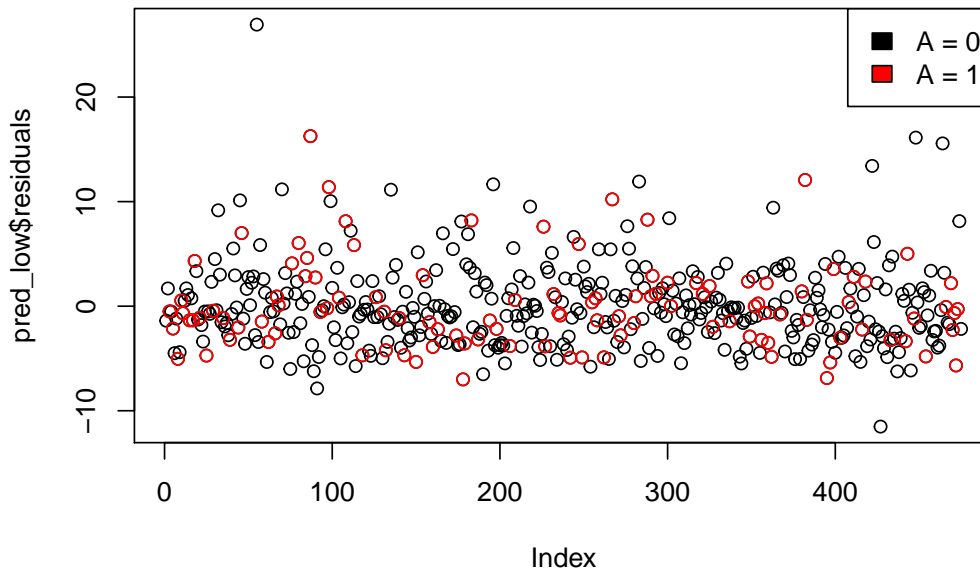
```
ATE_regression_adjustment_low = pred_low$coefficients[2]
ATE_regression_adjustment_low
```

```
##          A
## 3.05324
```

```
end.time_regression_adjustment_low <- Sys.time()
time_regression_adjustment_low <- end.time_regression_adjustment_low -
  start.time_regression_adjustment_low
time_regression_adjustment_low
```

```
## Time difference of 0.01595902 secs
```

### Residual Plot of Regression Adjustment Model – Low Dim



We find that the ATE for the low dimensional dataset was 3.0532404 with a runtime of 0.0159590244293213 seconds.

### Stratification and Regression Adjustment

In this last method, we will combine the first two methods together. In the same way as stratification, we split the datasets into  $K = 3$  strata. The choice of  $K = 3$  is again because it was the highest value of  $K$  that did not give us an empty stratum. Within each stratum, we perform regression adjustment by regressing

the response variable ( $Y$ ) with the treatment variable ( $A$ ) and the estimated propensity scores. We then have three coefficients for the variable  $A$ , one for each regression model. We then take a weighted average of these coefficients, with the weights relative to the population of each strata.

As D'Agostino (1998) notes, stratification combined with regression adjustment helps to reduce the bias in the treatment effect if the treatment groups are parallel and performs much better than propensity score matching alone. However, this method has the same drawbacks as stated in both stratification and regression adjustment.

```
##   A stratum   n   prop    avg_y
## 1 0         1   9 0.0045 -1.556754
## 2 1         1   1 0.0005  3.448809
## 3 0        2 805 0.4025 -13.679379
## 4 1        2 465 0.2325 -16.168680
## 5 0        3 289 0.1445 -6.706643
## 6 1        3 431 0.2155 -8.342732
```

```
##   A stratum   n      prop    avg_y
## 1 0         1  15 0.03157895 18.25654
## 2 1         1   0 0.00000000  0.00000
## 3 0         2 262 0.55157895 15.34302
## 4 1         2  34 0.07157895 18.68881
## 5 0         3  86 0.18105263 19.14737
## 6 1         3  78 0.16421053 22.52296
```

In particular, we see very imbalanced groups within each strata, as shown in the summary tables from the stratification method shown above, which may end up increasing the bias in our estimate of the ATE instead. We would be cautious of advocating for this method even if the ATE estimate was accurate.

```
start.time_stratification_regression_adjustment_high <- Sys.time()

lm_beta_high <- rep(NA, K)

for (i in 1:K){
  subset <- df_high[df_high$stratum_class_high == i, ]

  if (nrow(subset) == 0) {
    # if the stratum is empty, let the coefficient for A automatically be 0
    lm_beta_high[i] <- 0
  } else if (sum(subset$prop_score_high) == 0) {
    # if the propensity scores in the stratum are all 0,
    # let the coefficient for A automatically be 0
    lm_beta_low[i] <- 0
  } else {
    # otherwise, run a linear model on the subset
    lm <- lm(Y ~ A + prop_score_high, data = subset)
    lm_beta_high[i] <- as.numeric(lm$coefficients[2])
  }
}

lm_beta_high
```



## High Dimensional Data

```
## [1] 5.005563 -2.486942 -2.637645
```

```
ATE_stratification_regression_adjustment_high <- stratum_prop_high$sum[1]*lm_beta_high[1] +  
  stratum_prop_high$sum[2]*lm_beta_high[2] +  
  stratum_prop_high$sum[3]*lm_beta_high[3]  
  
ATE_stratification_regression_adjustment_high
```

```
## [1] -2.503732
```

```
end.time_stratification_regression_adjustment_high <- Sys.time()  
  
time_stratification_regression_adjustment_high <-  
  end.time_stratification_regression_adjustment_high -  
  start.time_stratification_regression_adjustment_high  
  
time_stratification_regression_adjustment_high
```

```
## Time difference of 0.03195 secs
```

We find that the ATE for the high dimensional dataset was -2.504 with a runtime of 0.032 seconds.

```
start.time_stratification_regression_adjustment_low <- Sys.time()  
  
lm_beta_low <- rep(NA, K)  
  
for (i in 1:K){  
  subset <- df_low[df_low$stratum_class_low == i, ]  
  
  if (nrow(subset) == 0) {  
    # if the stratum is empty, let the coefficient for A automatically be 0  
    lm_beta_low[i] <- 0  
  } else if (sum(subset$prop_score_low) == 0) {  
    # if the propensity scores in the stratum are all 0  
    # let the coefficient for A automatically be 0  
    lm_beta_low[i] <- 0  
  } else {  
    # otherwise, run a linear model on the subset  
    lm <- lm(Y ~ A + prop_score_low, data = subset)  
    lm_beta_low[i] <- as.numeric(lm$coefficients[2])  
  }  
}  
  
lm_beta_low
```

## Low Dimensional Data

```
## [1] 0.000000 3.264136 2.863805
```

```
ATE_stratification_regression_adjustment_low <- stratum_prop_low$sum[1]*lm_beta_low[1] +  
  stratum_prop_low$sum[2]*lm_beta_low[2] +  
  stratum_prop_low$sum[3]*lm_beta_low[3]
```

```
ATE_stratification_regression_adjustment_low
```

```
## [1] 3.022839
```

```
end.time_stratification_regression_adjustment_low <- Sys.time()
```

```
time_stratification_regression_adjustment_low <-  
  end.time_stratification_regression_adjustment_low -  
  start.time_stratification_regression_adjustment_low
```

```
time_stratification_regression_adjustment_low
```

```
## Time difference of 0.02493501 secs
```

We find that the ATE for the low dimensional dataset was 3.023 with a runtime of 0.025 seconds.

## Results

We compare the accuracy and performance of the three ATE Estimation procedures below.

### ATE Results

We are provided the true ATE values of -3 for the high dimensional data and 2.5 for the low dimensional data.

##	High Dimensional Data
## True	-3.000000
## Stratification	-2.144670
## Regression Adjustment	-2.527116
## Stratification + Regression Adjustment	-2.503732
##	Low Dimensional Data
## True	2.500000
## Stratification	2.673899
## Regression Adjustment	3.053240
## Stratification + Regression Adjustment	3.022839

From the table above, we see that regression adjustment performed the best for the high dimensional data and stratification performed the best for the low dimensional data.

### Runtime results

##	High Dimensional Data
## Propensity Score Estimation	1.24324393
## Stratification	0.31650686

## Regression Adjustment	0.04108310
## Stratification + Regression Adjustment	0.03195000
##	Low Dimensional Data
## Propensity Score Estimation	0.03291011
## Stratification	0.12155604
## Regression Adjustment	0.01595902
## Stratification + Regression Adjustment	0.02493501

Given the nature of trees, propensity score estimations are quickly calculated once we have the proper hyperparameters selected from cross-validation—even for the high dimensional data, propensity score estimations did not take more than two seconds.

It is also no surprise that, given the sizes of our two datasets, that regression adjustment was the fastest method. However, with larger datasets with more observations, this may not be the case. Stratification took the longest time, mainly due to the many intermediate calculations required. Lastly, the combination method of both stratification and regression adjustment had a runtime between the two former methods.

However, we want to note that this .Rmd file was knitted using a computer with a NVMe SAMSUNG SSD with 16 GB RAM. Runtimes may vary from device to device.

## Conclusion

Overall, we believe that using classification/regression trees for propensity scores was not the ideal approach for either dataset. While we cross-validated the complexity hyperparameter,  $cp$ , to help avoid with overfitting, our models for both the high dimensional and low dimensional datasets ended up estimating the same propensity score value for over half of the entire dataset. This would not be a very helpful model in differentiating our observations and of course affect our ATE estimations regardless of the method used.

We see this most prominently in stratification, in which different values of  $K$ , that is, the number of strata, resulted in an empty stratum in our results. Even after choosing a value of  $K$  which would present no empty strata, we saw that each stratum tend to have imbalanced classes. In the case of the low dimensional dataset, one stratum only consisted of observations from the control group. These complications may explain why the stratification plus regression adjustment method would not have performed the best.

However, the results were relatively consistent among all three methods—there were no large deviations from the true value. In particular, the ATE for stratification was actually quite close to the true value for the low dimensional data. Additionally, compared to other methods, we note the relative ease of interpretation and fast run times for not only the propensity score estimations but also for the ATE estimations as well. While we may not advocate for these estimation methods for their accuracy (and validity in certain cases), but these methods here show a fast and easy way to get a general sense of the average treatment effect.

## References

- Austin, Peter C. 2011. “An Introduction to Propensity Score Methods for Reducing the Effects of Confounding in Observational Studies.” *Multivariate Behavioral Research* 46 (3): 399–424.
- Chan, David & Ge, Rong & Gershony, Ori & Hesterberg, Tim & Lambert, Diane. (2010). Evaluating online ad campaigns in a pipeline: Causal models at scale. *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 7-16. 10.1145/1835804.1835809.
- D’Agostino RB Jr. Propensity score methods for bias reduction in the comparison of a treatment to a non-randomized control group. *Stat Med*. 1998 Oct 15;17(19):2265-81. doi: 10.1002/(sici)1097-0258(19981015)17:19<2265::aid-sim918>3.0.co;2-b. PMID: 9802183.

- Lunceford, Jared K, and Marie Davidian. 2004. “Stratification and Weighting via the Propensity Score in Estimation of Causal Treatment Effects a Comparative Study.” *Statistics in Medicine* 23 (19): 2937–60.